

CNN Avec peu de données

Chargement du code

```
In [1]: import numpy as np
from matplotlib import pyplot as plt
import tensorflow as tf

np.random.seed(123) # for reproducibility
tf.random.set_seed(123)

%load_ext autoreload
%autoreload 2
```

```
In [2]: from TP3_utils import load_dataset

DATA_PATH = 'Data'
CATEGORIES = ["accordion", "anchor", "barrel", "binocular"]

(X_train, Y_train), (X_test, Y_test), num_classes = load_dataset(DATA_PATH)
data_shape = X_train[0].shape

print("Training data shape:", data_shape)
print("Training labels shape:", Y_train.shape)
print("Test data shape:", X_test.shape)
print("Test labels shape:", Y_test.shape)
print("Number of classes:", num_classes)
```

Finished loading 177 images from 4 categories.

Train/Test split: 123 / 54

Training data shape: (224, 224, 3)

Training labels shape: (123, 4)

Test data shape: (54, 224, 224, 3)

Test labels shape: (54, 4)

Number of classes: 4

Finished loading 177 images from 4 categories.

Train/Test split: 123 / 54

Training data shape: (224, 224, 3)

Training labels shape: (123, 4)

Test data shape: (54, 224, 224, 3)

Test labels shape: (54, 4)

Number of classes: 4

Les images d'entrées sont redimensionnées à 224x224, et en couleurs (3 canaux pour RGB).

Il y a 4 classes en sortie, et la répartition est de 70% / 30% soit 123 exemples d'entraînement et 54 de test sur les 177 chargés.

```
In [ ]: import time
import tensorflow as tf
from tensorflow import keras
```

```

from keras.optimizers import Adam
from TP3_utils import CNN, affiche, eval_classif, data_augmentation
from keras.callbacks import ReduceLROnPlateau, EarlyStopping

lr_scheduler = ReduceLROnPlateau(monitor='val_loss', factor=0.75, patience=2)
earlStop = EarlyStopping(monitor='val_loss', min_delta=1e-4, patience=5)

lr=1e-4
batch_size=32
epochs=64
ad= Adam(learning_rate=lr)

model = CNN(data_shape)
model.summary()

plt.figure(figsize=(10, 4))
for i in range(5):
    augmented_img = data_augmentation(tf.expand_dims(X_train[0], 0))
    plt.subplot(1, 5, i+1)
    plt.imshow(augmented_img[0])
    plt.axis("off")
    plt.show()

model.compile(
    loss='categorical_crossentropy',
    optimizer=ad,
    metrics=['accuracy']
)

tps1 = time.time()
history = model.fit(
    X_train,
    Y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(X_test, Y_test),
    callbacks=[earlStop, lr_scheduler]
)
tps2 = time.time()

# Evaluation
loss, accuracy = model.evaluate(X_test, Y_test)
print(f'Test loss: {loss}, Test accuracy: {accuracy}')

affiche(history)
preds = model.predict(X_test)
eval_classif(Y_test, preds)
print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))

```

Model: "functional_1"

Layer (type)	Output Shape
input_layer (InputLayer)	(None, 224, 224, 3)
sequential (Sequential)	(None, 224, 224, 3)
conv2d (Conv2D)	(None, 224, 224, 16)
max_pooling2d (MaxPooling2D)	(None, 112, 112, 16)
conv2d_1 (Conv2D)	(None, 112, 112, 32)
conv2d_2 (Conv2D)	(None, 112, 112, 32)
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 32)
dropout (Dropout)	(None, 56, 56, 32)
conv2d_3 (Conv2D)	(None, 56, 56, 64)
conv2d_4 (Conv2D)	(None, 56, 56, 64)
dropout_1 (Dropout)	(None, 56, 56, 64)
flatten (Flatten)	(None, 200704)
dense (Dense)	(None, 100)
dropout_2 (Dropout)	(None, 100)
dense_1 (Dense)	(None, 4)

Total params: 20,140,664 (76.83 MB)

Trainable params: 20,140,664 (76.83 MB)

Non-trainable params: 0 (0.00 B)





```
Epoch 1/64
4/4 3s 497ms/step - accuracy: 0.3089 - loss: 1.
4556 - val_accuracy: 0.4815 - val_loss: 1.3045 - learning_rate: 1.00
00e-04
Epoch 2/64
4/4 2s 445ms/step - accuracy: 0.3415 - loss: 1.
3691 - val_accuracy: 0.6667 - val_loss: 1.1482 - learning_rate: 1.00
00e-04
Epoch 3/64
4/4 2s 432ms/step - accuracy: 0.4309 - loss: 1.
2772 - val_accuracy: 0.6667 - val_loss: 1.0795 - learning_rate: 1.00
00e-04
Epoch 4/64
4/4 2s 427ms/step - accuracy: 0.4634 - loss: 1.
2088 - val_accuracy: 0.7593 - val_loss: 0.9737 - learning_rate: 1.00
00e-04
Epoch 5/64
4/4 2s 426ms/step - accuracy: 0.5041 - loss: 1.
1355 - val_accuracy: 0.7778 - val_loss: 0.8628 - learning_rate: 1.00
00e-04
Epoch 6/64
4/4 2s 427ms/step - accuracy: 0.5854 - loss: 1.
0419 - val_accuracy: 0.7222 - val_loss: 0.7416 - learning_rate: 1.00
00e-04
Epoch 7/64
4/4 2s 428ms/step - accuracy: 0.6423 - loss: 1.
0069 - val_accuracy: 0.7778 - val_loss: 0.6771 - learning_rate: 1.00
00e-04
Epoch 8/64
4/4 2s 429ms/step - accuracy: 0.6179 - loss: 0.
9297 - val_accuracy: 0.7407 - val_loss: 0.6265 - learning_rate: 1.00
00e-04
Epoch 9/64
4/4 2s 426ms/step - accuracy: 0.6179 - loss: 0.
9272 - val_accuracy: 0.7963 - val_loss: 0.6031 - learning_rate: 1.00
00e-04
Epoch 10/64
```

4/4 2s 426ms/step - accuracy: 0.6911 - loss: 0.8128 - val_accuracy: 0.7222 - val_loss: 0.5975 - learning_rate: 1.00e-04
Epoch 11/64
4/4 2s 426ms/step - accuracy: 0.6504 - loss: 0.7754 - val_accuracy: 0.7593 - val_loss: 0.5500 - learning_rate: 1.00e-04
Epoch 12/64
4/4 2s 438ms/step - accuracy: 0.6585 - loss: 0.7587 - val_accuracy: 0.7407 - val_loss: 0.6055 - learning_rate: 1.00e-04
Epoch 13/64
4/4 2s 429ms/step - accuracy: 0.6504 - loss: 0.7779 - val_accuracy: 0.7593 - val_loss: 0.5290 - learning_rate: 1.00e-04
Epoch 14/64
4/4 2s 438ms/step - accuracy: 0.6992 - loss: 0.6915 - val_accuracy: 0.7407 - val_loss: 0.5553 - learning_rate: 1.00e-04
Epoch 15/64
4/4 0s 370ms/step - accuracy: 0.6880 - loss: 0.7282
Epoch 15: ReduceLROnPlateau reducing learning rate to 7.499999810534064e-05.
4/4 2s 428ms/step - accuracy: 0.7154 - loss: 0.6690 - val_accuracy: 0.7963 - val_loss: 0.5338 - learning_rate: 1.00e-04
Epoch 16/64
4/4 2s 428ms/step - accuracy: 0.7154 - loss: 0.6758 - val_accuracy: 0.7963 - val_loss: 0.5221 - learning_rate: 7.50e-05
Epoch 17/64
4/4 2s 427ms/step - accuracy: 0.7805 - loss: 0.6097 - val_accuracy: 0.7778 - val_loss: 0.5461 - learning_rate: 7.50e-05
Epoch 18/64
4/4 2s 429ms/step - accuracy: 0.7480 - loss: 0.6374 - val_accuracy: 0.7778 - val_loss: 0.4941 - learning_rate: 7.50e-05
Epoch 19/64
4/4 2s 426ms/step - accuracy: 0.7724 - loss: 0.5386 - val_accuracy: 0.8148 - val_loss: 0.5046 - learning_rate: 7.50e-05
Epoch 20/64
4/4 0s 372ms/step - accuracy: 0.7925 - loss: 0.5780
Epoch 20: ReduceLROnPlateau reducing learning rate to 5.6249997214763425e-05.
4/4 2s 430ms/step - accuracy: 0.8211 - loss: 0.5295 - val_accuracy: 0.7778 - val_loss: 0.5220 - learning_rate: 7.50e-05
Epoch 21/64
4/4 2s 427ms/step - accuracy: 0.8130 - loss: 0.4985 - val_accuracy: 0.8333 - val_loss: 0.4722 - learning_rate: 5.6250e-05
Epoch 22/64

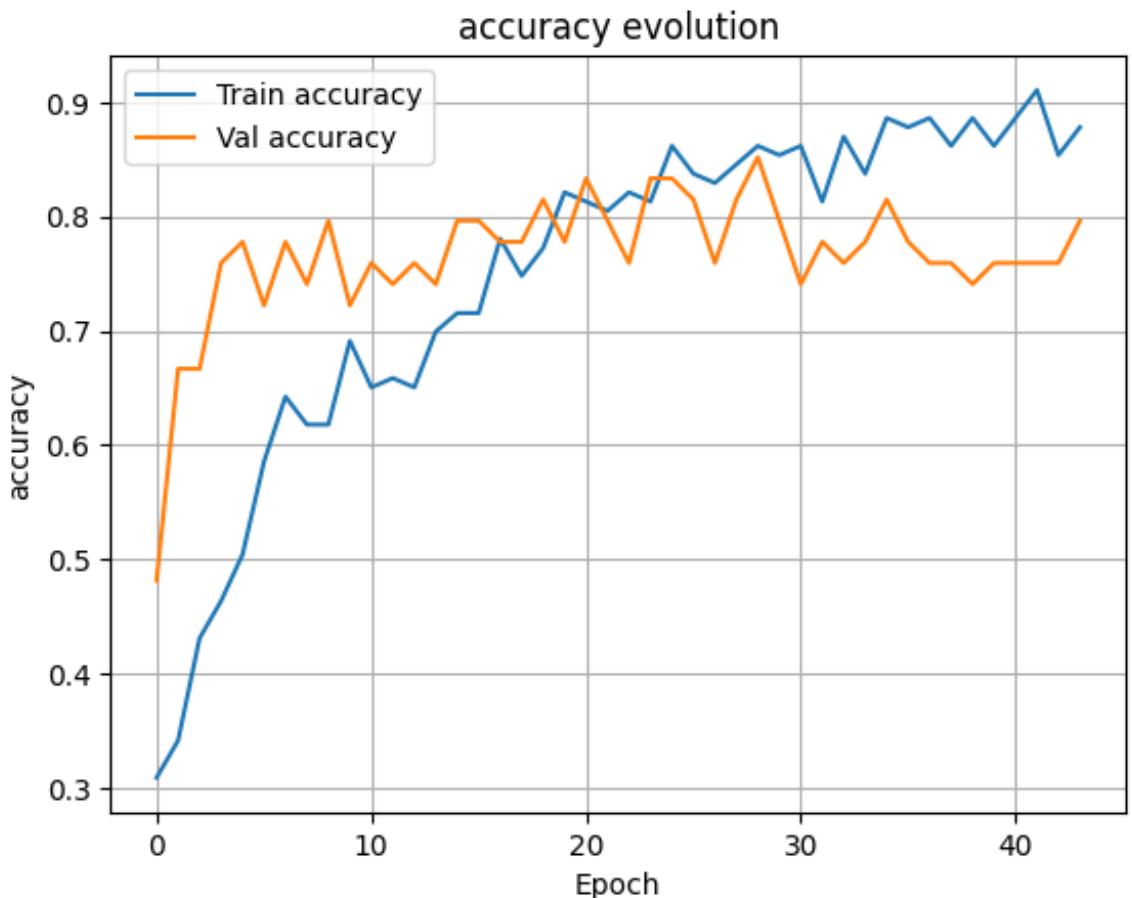
4/4 2s 427ms/step - accuracy: 0.8049 - loss: 0.
5010 - val_accuracy: 0.7963 - val_loss: 0.4760 - learning_rate: 5.62
50e-05
Epoch 23/64
4/4 0s 379ms/step - accuracy: 0.8212 - loss: 0.
4520
Epoch 23: ReduceLROnPlateau reducing learning rate to 4.218749927531
462e-05.
4/4 2s 437ms/step - accuracy: 0.8211 - loss: 0.
4596 - val_accuracy: 0.7593 - val_loss: 0.5322 - learning_rate: 5.62
50e-05
Epoch 24/64
4/4 2s 431ms/step - accuracy: 0.8130 - loss: 0.
4697 - val_accuracy: 0.8333 - val_loss: 0.4646 - learning_rate: 4.21
88e-05
Epoch 25/64
4/4 2s 426ms/step - accuracy: 0.8618 - loss: 0.
4183 - val_accuracy: 0.8333 - val_loss: 0.4563 - learning_rate: 4.21
88e-05
Epoch 26/64
4/4 2s 428ms/step - accuracy: 0.8374 - loss: 0.
4422 - val_accuracy: 0.8148 - val_loss: 0.4702 - learning_rate: 4.21
88e-05
Epoch 27/64
4/4 0s 370ms/step - accuracy: 0.8219 - loss: 0.
3739
Epoch 27: ReduceLROnPlateau reducing learning rate to 3.164062582072
802e-05.
4/4 2s 428ms/step - accuracy: 0.8293 - loss: 0.
4005 - val_accuracy: 0.7593 - val_loss: 0.4711 - learning_rate: 4.21
88e-05
Epoch 28/64
4/4 2s 428ms/step - accuracy: 0.8455 - loss: 0.
4143 - val_accuracy: 0.8148 - val_loss: 0.4388 - learning_rate: 3.16
41e-05
Epoch 29/64
4/4 2s 431ms/step - accuracy: 0.8618 - loss: 0.
3893 - val_accuracy: 0.8519 - val_loss: 0.4269 - learning_rate: 3.16
41e-05
Epoch 30/64
4/4 2s 426ms/step - accuracy: 0.8537 - loss: 0.
3806 - val_accuracy: 0.7963 - val_loss: 0.4611 - learning_rate: 3.16
41e-05
Epoch 31/64
4/4 0s 369ms/step - accuracy: 0.8274 - loss: 0.
3692
Epoch 31: ReduceLROnPlateau reducing learning rate to 2.373046936554
6014e-05.
4/4 2s 427ms/step - accuracy: 0.8618 - loss: 0.
3584 - val_accuracy: 0.7407 - val_loss: 0.4511 - learning_rate: 3.16
41e-05
Epoch 32/64
4/4 2s 431ms/step - accuracy: 0.8130 - loss: 0.
4602 - val_accuracy: 0.7778 - val_loss: 0.4614 - learning_rate: 2.37
30e-05
Epoch 33/64

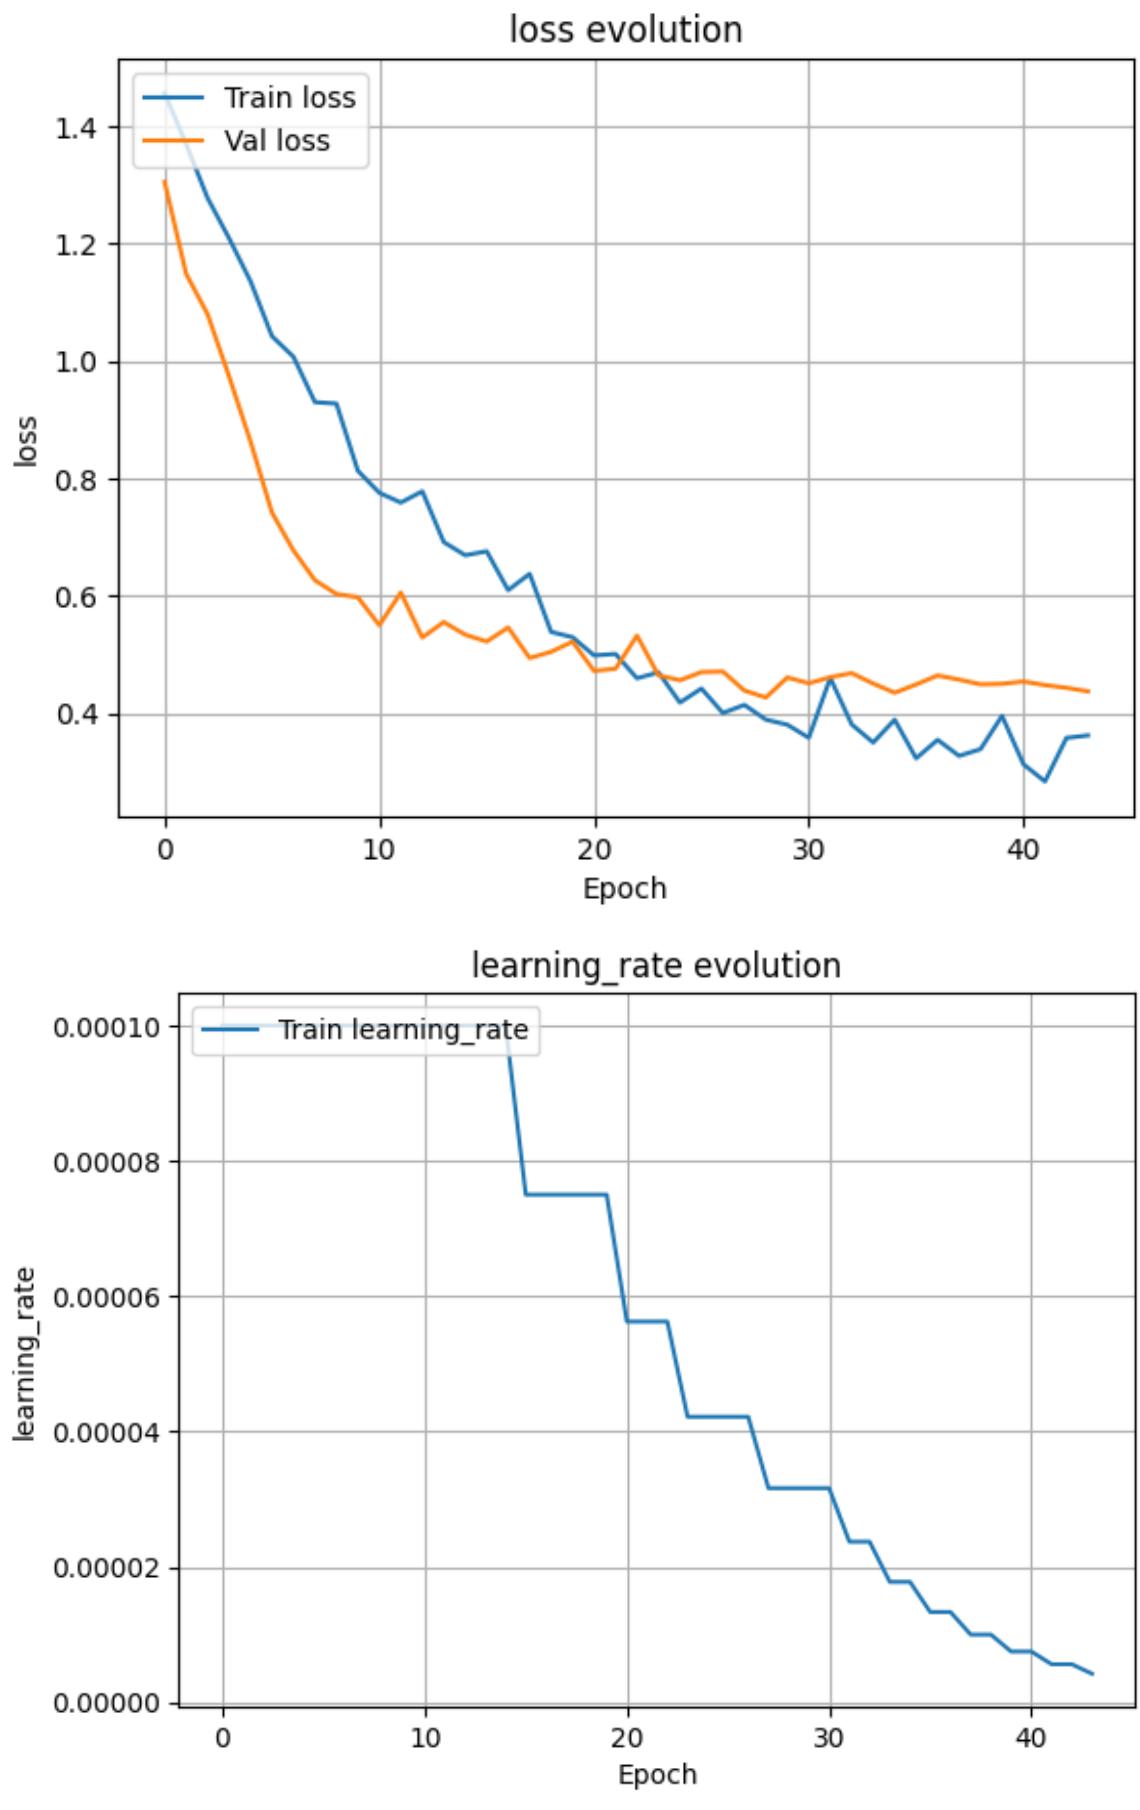
4/4 ━━━━━━ 0s 371ms/step - accuracy: 0.8998 - loss: 0.3796
Epoch 33: ReduceLROnPlateau reducing learning rate to 1.7797852706280537e-05.
4/4 ━━━━━━ 2s 438ms/step - accuracy: 0.8699 - loss: 0.3810 - val_accuracy: 0.7593 - val_loss: 0.4684 - learning_rate: 2.3730e-05
Epoch 34/64
4/4 ━━━━━━ 2s 428ms/step - accuracy: 0.8374 - loss: 0.3503 - val_accuracy: 0.7778 - val_loss: 0.4507 - learning_rate: 1.7798e-05
Epoch 35/64
4/4 ━━━━━━ 0s 373ms/step - accuracy: 0.8986 - loss: 0.4195
Epoch 35: ReduceLROnPlateau reducing learning rate to 1.3348389529710403e-05.
4/4 ━━━━━━ 2s 431ms/step - accuracy: 0.8862 - loss: 0.3891 - val_accuracy: 0.8148 - val_loss: 0.4350 - learning_rate: 1.7798e-05
Epoch 36/64
4/4 ━━━━━━ 2s 428ms/step - accuracy: 0.8780 - loss: 0.3234 - val_accuracy: 0.7778 - val_loss: 0.4491 - learning_rate: 1.3348e-05
Epoch 37/64
4/4 ━━━━━━ 0s 371ms/step - accuracy: 0.8921 - loss: 0.3341
Epoch 37: ReduceLROnPlateau reducing learning rate to 1.0011292488343315e-05.
4/4 ━━━━━━ 2s 429ms/step - accuracy: 0.8862 - loss: 0.3549 - val_accuracy: 0.7593 - val_loss: 0.4645 - learning_rate: 1.3348e-05
Epoch 38/64
4/4 ━━━━━━ 2s 430ms/step - accuracy: 0.8618 - loss: 0.3275 - val_accuracy: 0.7593 - val_loss: 0.4574 - learning_rate: 1.0011e-05
Epoch 39/64
4/4 ━━━━━━ 0s 370ms/step - accuracy: 0.8986 - loss: 0.3300
Epoch 39: ReduceLROnPlateau reducing learning rate to 7.508469025196973e-06.
4/4 ━━━━━━ 2s 429ms/step - accuracy: 0.8862 - loss: 0.3390 - val_accuracy: 0.7407 - val_loss: 0.4494 - learning_rate: 1.0011e-05
Epoch 40/64
4/4 ━━━━━━ 2s 429ms/step - accuracy: 0.8618 - loss: 0.3956 - val_accuracy: 0.7593 - val_loss: 0.4502 - learning_rate: 7.5085e-06
Epoch 41/64
4/4 ━━━━━━ 0s 370ms/step - accuracy: 0.8713 - loss: 0.3323
Epoch 41: ReduceLROnPlateau reducing learning rate to 5.63135176889773e-06.
4/4 ━━━━━━ 2s 428ms/step - accuracy: 0.8862 - loss: 0.3132 - val_accuracy: 0.7593 - val_loss: 0.4542 - learning_rate: 7.5085e-06
Epoch 42/64

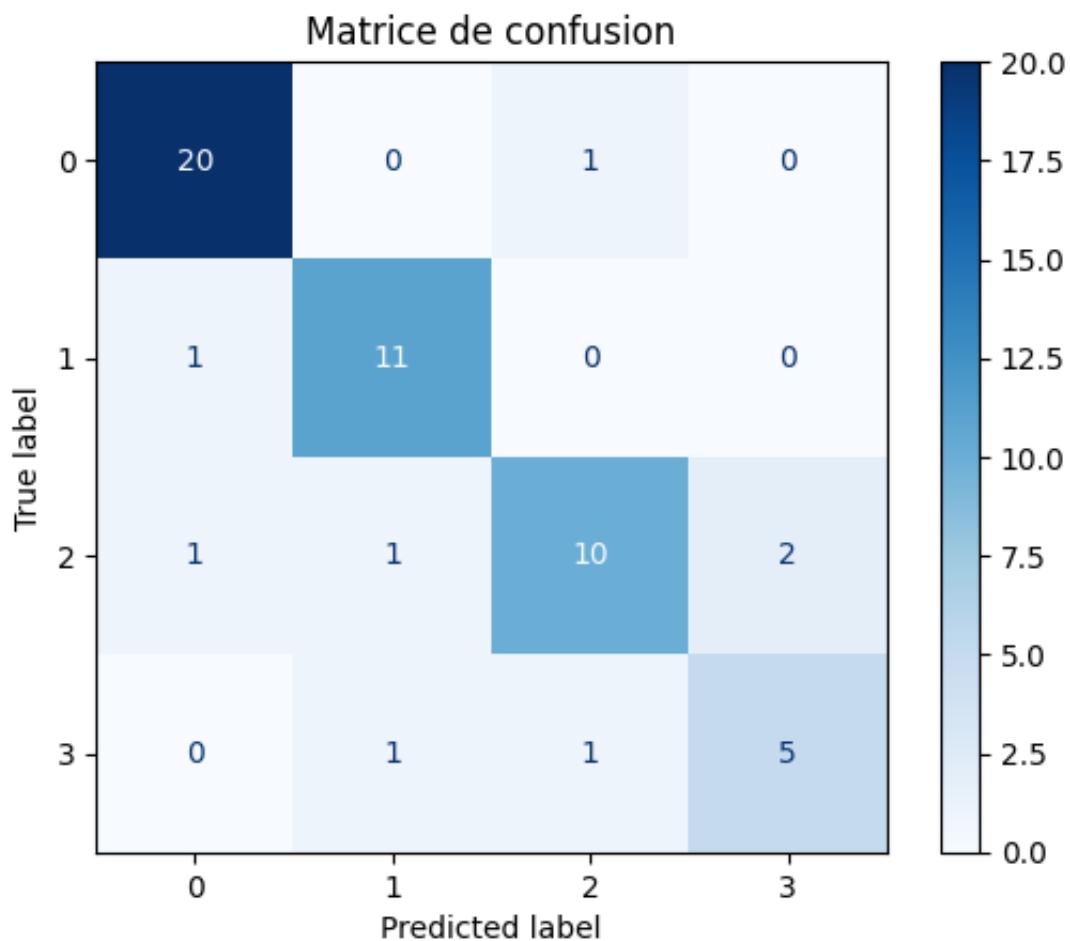
```

4/4 ━━━━━━━━ 2s 428ms/step - accuracy: 0.9106 - loss: 0.
2840 - val_accuracy: 0.7593 - val_loss: 0.4480 - learning_rate: 5.63
14e-06
Epoch 43/64
4/4 ━━━━━━━━ 0s 370ms/step - accuracy: 0.8111 - loss: 0.
4033
Epoch 43: ReduceLROnPlateau reducing learning rate to 4.223513997203
554e-06.
4/4 ━━━━━━━━ 2s 438ms/step - accuracy: 0.8537 - loss: 0.
3581 - val_accuracy: 0.7593 - val_loss: 0.4434 - learning_rate: 5.63
14e-06
Epoch 44/64
4/4 ━━━━━━━━ 2s 430ms/step - accuracy: 0.8780 - loss: 0.
3624 - val_accuracy: 0.7963 - val_loss: 0.4374 - learning_rate: 4.22
35e-06
Epoch 44: early stopping
Restoring model weights from the end of the best epoch: 29.
2/2 ━━━━━━━━ 0s 73ms/step - accuracy: 0.8519 - loss: 0.4
269
Test loss: 0.42687469720840454, Test accuracy: 0.8518518805503845

```







Classification report:

	precision	recall	f1-score	support
0	0.91	0.95	0.93	21
1	0.85	0.92	0.88	12
2	0.83	0.71	0.77	14
3	0.71	0.71	0.71	7
accuracy			0.85	54
macro avg	0.83	0.82	0.82	54
weighted avg	0.85	0.85	0.85	54

Temps d'entraînement : 75.47 secondes

Le modèle initial à **20,090,296** paramètres entraînable. Après un entraînement de **19.32s**, le modèle atteint **loss=0.8; accuracy=66.66%**

Pour définir la structure du modèle, je fixe:

- lr=1e-4
- batch_size=16
- epochs=16
- optimizer=Adam(lr)
- MLP finale: Dense(100)

Puis, je procède de la façon suivante:

itération	nb étages	taille des filtres	nb params	tps d'entraînement	loss	accuracy
1	2	16-32	40,146,392	23.37s	0.7255	72.22%
2	3	16-32-64	20,094,488	19.03s	0.5276	77.77%
3	4	16-32-64-128	10,133,144	18.83s	0.8526	72.22%
4	5	16-32-64-128-256	5,410,712	20.09s	0.9057	70.37%
5	3	32-64-128	40,234,552	47.86s	0.8762	72.22%
6	4	32-64-128-256	20,459,320	44.29s	0.9546	75.92%
7	5	32-64-128-256-512	11,604,280	51.48s	0.7749	74.07%
8	2	8-16	20,072,296	13.89s	0.7267	66.66%
9	3	8-16-32	10,041,736	12.12s	0.7994	68.51%

Pour les N premiers, overfitting très rapide; aucun réel apprentissage après quelques epochs seulement. La majorité des paramètres sont à la couche dense après le flatten, donc augmenter le nombre de couche permet de faire intervenir plus de max pooling, donc réduire le temps d'apprentissage tout en ayant une meilleure répartition des paramètres le long du réseau. Cependant, un réseau trop long réduit trop le nombre de paramètres, donc la capacité d'expression du réseau? Ou bien est-ce un souci d'évanescence du gradient? On peut alors tenter d'augmenter le nombre de filtres, mais on remarque que cela n'améliore pas la précision et ralentit largement l'apprentissage. On constate nettement des signes d'overfitting.

Au contraire, en tentant de réduire le nombre de filtre, on réduit le temps d'apprentissage mais là encore, la précision diminue.

Le meilleur compromis semble alors être d'utiliser 3 étages, avec respectivement 16, 32 et 64 filtres. Pour enrichir le réseau, on peut maintenant essayer d'augmenter le nombre de couches par étages.

itération	nb couches/ étage	nb params	tps d'entraînement (total)	loss	best val_accuracy
0	1-1-1	20,094,488	19.03s	0.5276	77.77%
1	2-1-1	20,096,808	28.84s	0.8320	75.92%
2	1-2-1	20,103,736	25.42s	0.9896	72.22%
3	1-1-2	20,131,416	23.72s	0.9335	64.81%
4	2-2-1	20,106,056	34.36s	0.5813	74.07%
5	2-1-2	20,133,736	34.46s	0.6585	72.22%
6	1-2-2	20,140,664	31.39s	0.7302	79.62%
7	2-2-2	20,142,984	40.51s	0.7089	75.92%

Bien que cette fois, l'impact est moins clair, la configuration qui semble la meilleure est 1-2-2 (bien que 79% semble avoir été un outlier, cette configuration donne de façon plutôt consistante >73% ce qui n'est pas le cas des autres). Ce que j'utiliserai par la suite. Cependant on remarque un fort overfitting. Pour palier cela, nous ajouteraons du dropout pour les différentes layers, puis de la data augmentation.

La transformation qui me semble la plus efficace ici est une rotation légère. Aussi, avec data augmentation et dropout, l'entraînement était très instable, alors j'ai augmenté le batch_size. Malgré tout, le modèle était très sensible et je n'ai pas pu mettre beaucoup de dropout, surtout sur les couches hautes. En m'aidant d'early stopping et d'un lr scheduler, j'ai atteint au mieux ces résultats:

```
Classification report:
      precision    recall    f1-score   support
      0          0.91      0.95      0.93       21
      1          0.85      0.92      0.88       12
      2          0.83      0.71      0.77       14
      3          0.71      0.71      0.71        7

  accuracy                           0.85      54
  macro avg       0.83      0.82      0.82      54
weighted avg     0.85      0.85      0.85      54
```

Temps d'entraînement : 75.47 secondes

```
In [9]: from tensorflow.keras.layers import Input, Flatten
from tensorflow.keras.models import Model
from TP3_utils import MLP_transfer

lr=1e-4
batch_size=16
```

```

epochs=16
ad=Adam(learning_rate=lr)

# Transfer learning
input_tensor = Input(shape=(224,224,3))
VGG = tf.keras.applications.VGG16(
    weights='imagenet',
    include_top=False,
    input_tensor=input_tensor
)
for layer in VGG.layers:
    layer.trainable = False

x = VGG.output
x = Flatten()(x)
x = MLP_transfer(x, num_classes)
model_transfer = Model(inputs=VGG.input, outputs=x)

model_transfer.summary()

model_transfer.compile(
    loss='categorical_crossentropy',
    optimizer=Adam(1e-5),
    metrics=['accuracy']
)

tps1 = time.time()
history =model_transfer.fit(
    X_train,
    Y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(X_test, Y_test)
#    callbacks=[earlyStopBis, lr_schedulerBis],
)
tps2 = time.time()

loss, accuracy = model_transfer.evaluate(X_test, Y_test)
print(f'Test loss: {loss}, Test accuracy: {accuracy}')

affiche(history)
preds = model_transfer.predict(X_test)
eval_classif(Y_test, preds)
print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))

```

Model: "functional_8"

Layer (type)	Output Shape
input_layer_13 (InputLayer)	(None, 224, 224, 3)
block1_conv1 (Conv2D)	(None, 224, 224, 64)
block1_conv2 (Conv2D)	(None, 224, 224, 64)
block1_pool (MaxPooling2D)	(None, 112, 112, 64)
block2_conv1 (Conv2D)	(None, 112, 112, 128)
block2_conv2 (Conv2D)	(None, 112, 112, 128)
block2_pool (MaxPooling2D)	(None, 56, 56, 128)
block3_conv1 (Conv2D)	(None, 56, 56, 256)
block3_conv2 (Conv2D)	(None, 56, 56, 256)
block3_conv3 (Conv2D)	(None, 56, 56, 256)
block3_pool (MaxPooling2D)	(None, 28, 28, 256)
block4_conv1 (Conv2D)	(None, 28, 28, 512)
block4_conv2 (Conv2D)	(None, 28, 28, 512)
block4_conv3 (Conv2D)	(None, 28, 28, 512)
block4_pool (MaxPooling2D)	(None, 14, 14, 512)
block5_conv1 (Conv2D)	(None, 14, 14, 512)
block5_conv2 (Conv2D)	(None, 14, 14, 512)
block5_conv3 (Conv2D)	(None, 14, 14, 512)
block5_pool (MaxPooling2D)	(None, 7, 7, 512)
flatten_8 (Flatten)	(None, 25088)
dense_11 (Dense)	(None, 256)
dropout_3 (Dropout)	(None, 256)
dense_12 (Dense)	(None, 128)
dropout_4 (Dropout)	(None, 128)
dense_13 (Dense)	(None, 4)

Total params: 21,170,884 (80.76 MB)

Trainable params: 6,456,196 (24.63 MB)

Non-trainable params: 14,714,688 (56.13 MB)

Epoch 1/16
8/8 16s 2s/step - accuracy: 0.2927 - loss: 1.52
28 - val_accuracy: 0.3333 - val_loss: 1.3289

Epoch 2/16
8/8 15s 2s/step - accuracy: 0.3333 - loss: 1.44
40 - val_accuracy: 0.5185 - val_loss: 1.2417

Epoch 3/16
8/8 14s 2s/step - accuracy: 0.4065 - loss: 1.39
28 - val_accuracy: 0.6667 - val_loss: 1.1546

Epoch 4/16
8/8 15s 2s/step - accuracy: 0.3252 - loss: 1.37
78 - val_accuracy: 0.7222 - val_loss: 1.0697

Epoch 5/16
8/8 16s 2s/step - accuracy: 0.4472 - loss: 1.20
05 - val_accuracy: 0.7593 - val_loss: 0.9936

Epoch 6/16
8/8 15s 2s/step - accuracy: 0.4634 - loss: 1.27
04 - val_accuracy: 0.7963 - val_loss: 0.9298

Epoch 7/16
8/8 16s 2s/step - accuracy: 0.5122 - loss: 1.05
07 - val_accuracy: 0.8333 - val_loss: 0.8760

Epoch 8/16
8/8 16s 2s/step - accuracy: 0.6423 - loss: 0.98
67 - val_accuracy: 0.9259 - val_loss: 0.8171

Epoch 9/16
8/8 16s 2s/step - accuracy: 0.7073 - loss: 0.92
01 - val_accuracy: 0.9259 - val_loss: 0.7690

Epoch 10/16
8/8 16s 2s/step - accuracy: 0.6504 - loss: 0.88
47 - val_accuracy: 0.9259 - val_loss: 0.7190

Epoch 11/16
8/8 15s 2s/step - accuracy: 0.7154 - loss: 0.80
57 - val_accuracy: 0.9259 - val_loss: 0.6689

Epoch 12/16
8/8 15s 2s/step - accuracy: 0.7236 - loss: 0.75
59 - val_accuracy: 0.9259 - val_loss: 0.6255

Epoch 13/16
8/8 15s 2s/step - accuracy: 0.7724 - loss: 0.74
39 - val_accuracy: 0.9444 - val_loss: 0.5841

Epoch 14/16
8/8 15s 2s/step - accuracy: 0.8211 - loss: 0.65
76 - val_accuracy: 0.9444 - val_loss: 0.5339

Epoch 15/16
8/8 15s 2s/step - accuracy: 0.8537 - loss: 0.55
12 - val_accuracy: 0.9444 - val_loss: 0.4927

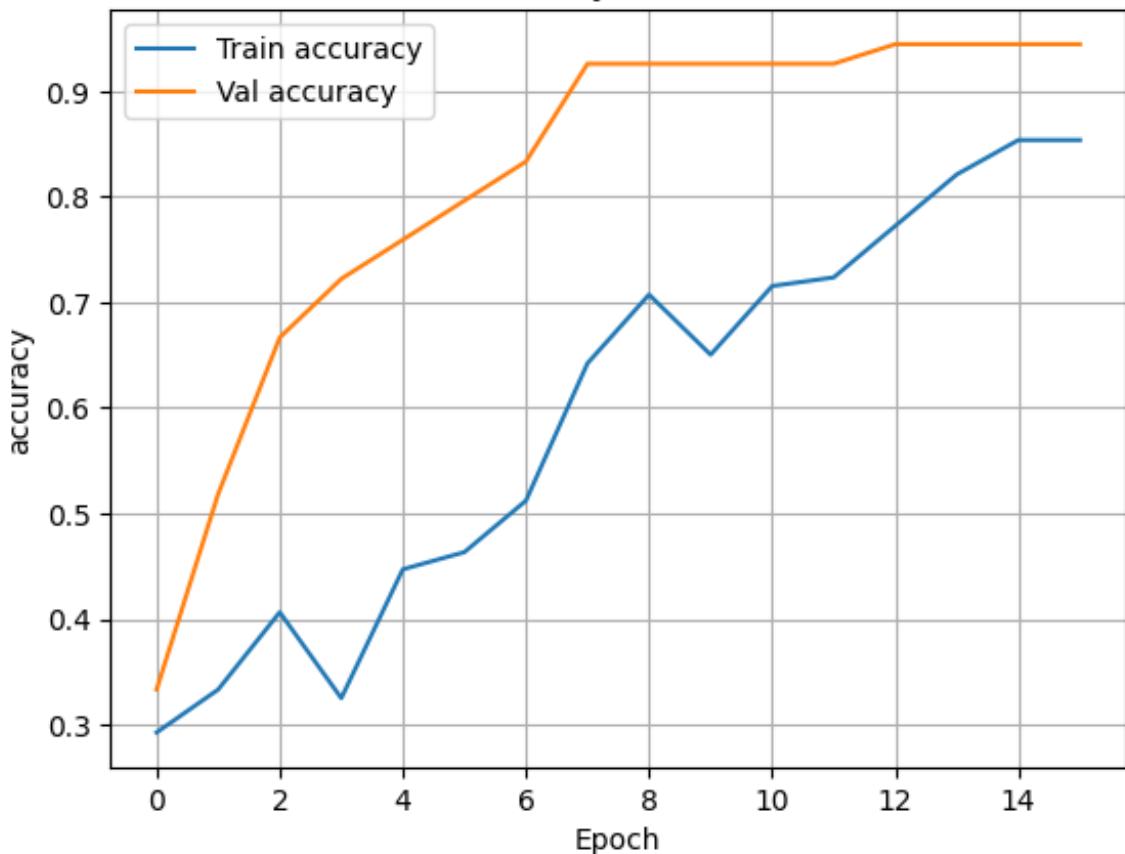
Epoch 16/16
8/8 15s 2s/step - accuracy: 0.8537 - loss: 0.56
17 - val_accuracy: 0.9444 - val_loss: 0.4633

2/2 5s 2s/step - accuracy: 0.9444 - loss: 0.463

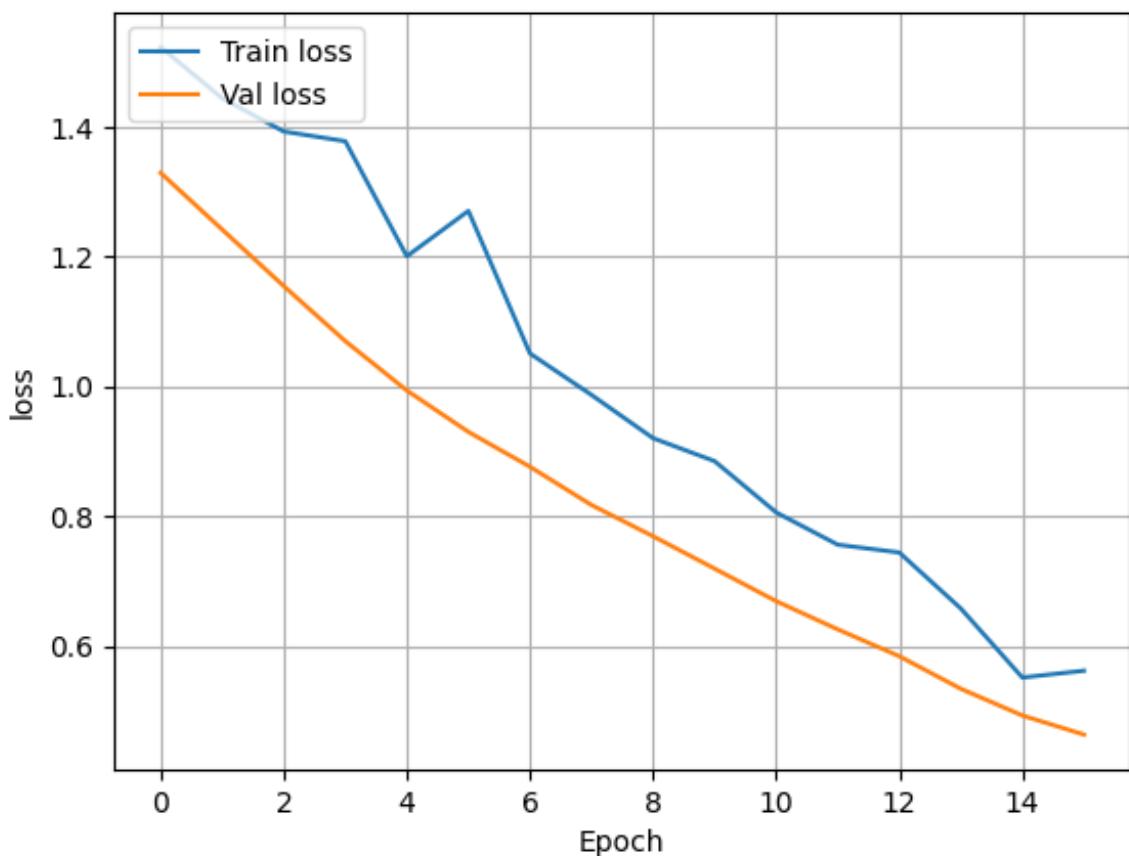
3

Test loss: 0.46325498819351196, Test accuracy: 0.9444444179534912

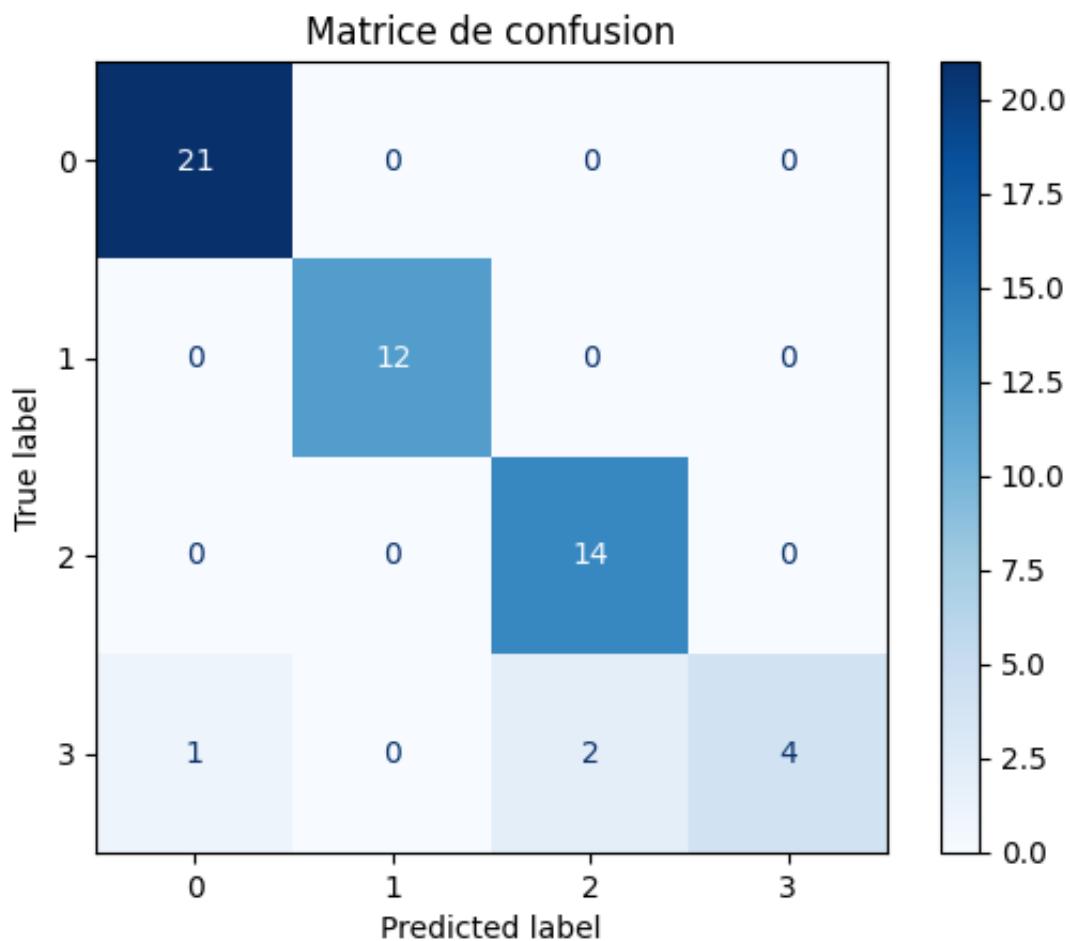
accuracy evolution



loss evolution



2/2 ————— 6s 2s/step



Classification report:

	precision	recall	f1-score	support
0	0.95	1.00	0.98	21
1	1.00	1.00	1.00	12
2	0.88	1.00	0.93	14
3	1.00	0.57	0.73	7
accuracy			0.94	54
macro avg	0.96	0.89	0.91	54
weighted avg	0.95	0.94	0.94	54

Temps d'entraînement : 243.76 secondes

On obtiens très facilement des performances nettement meilleures qu'avant:

Classification report:

	precision	recall	f1-score	support
0	0.95	1.00	0.98	21
1	1.00	1.00	1.00	12
2	0.88	1.00	0.93	14
3	1.00	0.57	0.73	7
accuracy			0.94	54
macro avg	0.96	0.89	0.91	54
weighted avg	0.95	0.94	0.94	54

Temps d'entraînement : 243.76 secondes

```
In [10]: ##### Fine-tuning #####

```

```
for layer in VGG.layers[-4:]:
    layer.trainable = True

model_transfer.compile(
    loss='categorical_crossentropy',
    optimizer=Adam(1e-6),
    metrics=['accuracy']
)

tps1 = time.time()
history = model_transfer.fit(
    X_train,
    Y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(X_test, Y_test)
#    callbacks=[earlyStopBis, lr_schedulerBis],
)
tps2 = time.time()

loss, accuracy = model_transfer.evaluate(X_test, Y_test)
print(f'Test loss: {loss}, Test accuracy: {accuracy}')

affiche(history)
preds = model_transfer.predict(X_test)
eval_classif(Y_test, preds)
print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))

```

Epoch 1/16
8/8 18s 2s/step - accuracy: 0.9024 - loss: 0.52
67 - val_accuracy: 0.9444 - val_loss: 0.4488

Epoch 2/16
8/8 17s 2s/step - accuracy: 0.8211 - loss: 0.54
06 - val_accuracy: 0.9444 - val_loss: 0.4355

Epoch 3/16
8/8 17s 2s/step - accuracy: 0.8537 - loss: 0.52
06 - val_accuracy: 0.9630 - val_loss: 0.4219

Epoch 4/16
8/8 17s 2s/step - accuracy: 0.8374 - loss: 0.47
91 - val_accuracy: 0.9630 - val_loss: 0.4101

Epoch 5/16
8/8 16s 2s/step - accuracy: 0.8455 - loss: 0.46
32 - val_accuracy: 0.9630 - val_loss: 0.3985

Epoch 6/16
8/8 17s 2s/step - accuracy: 0.8618 - loss: 0.49
10 - val_accuracy: 0.9630 - val_loss: 0.3869

Epoch 7/16
8/8 17s 2s/step - accuracy: 0.8455 - loss: 0.50
02 - val_accuracy: 0.9630 - val_loss: 0.3766

Epoch 8/16
8/8 16s 2s/step - accuracy: 0.8618 - loss: 0.50
68 - val_accuracy: 0.9630 - val_loss: 0.3664

Epoch 9/16
8/8 16s 2s/step - accuracy: 0.8862 - loss: 0.43
06 - val_accuracy: 0.9630 - val_loss: 0.3564

Epoch 10/16
8/8 16s 2s/step - accuracy: 0.9024 - loss: 0.44
19 - val_accuracy: 0.9630 - val_loss: 0.3469

Epoch 11/16
8/8 16s 2s/step - accuracy: 0.9268 - loss: 0.39
78 - val_accuracy: 0.9630 - val_loss: 0.3379

Epoch 12/16
8/8 16s 2s/step - accuracy: 0.9024 - loss: 0.41
33 - val_accuracy: 0.9630 - val_loss: 0.3296

Epoch 13/16
8/8 16s 2s/step - accuracy: 0.8943 - loss: 0.39
38 - val_accuracy: 0.9630 - val_loss: 0.3215

Epoch 14/16
8/8 16s 2s/step - accuracy: 0.9268 - loss: 0.36
30 - val_accuracy: 0.9630 - val_loss: 0.3129

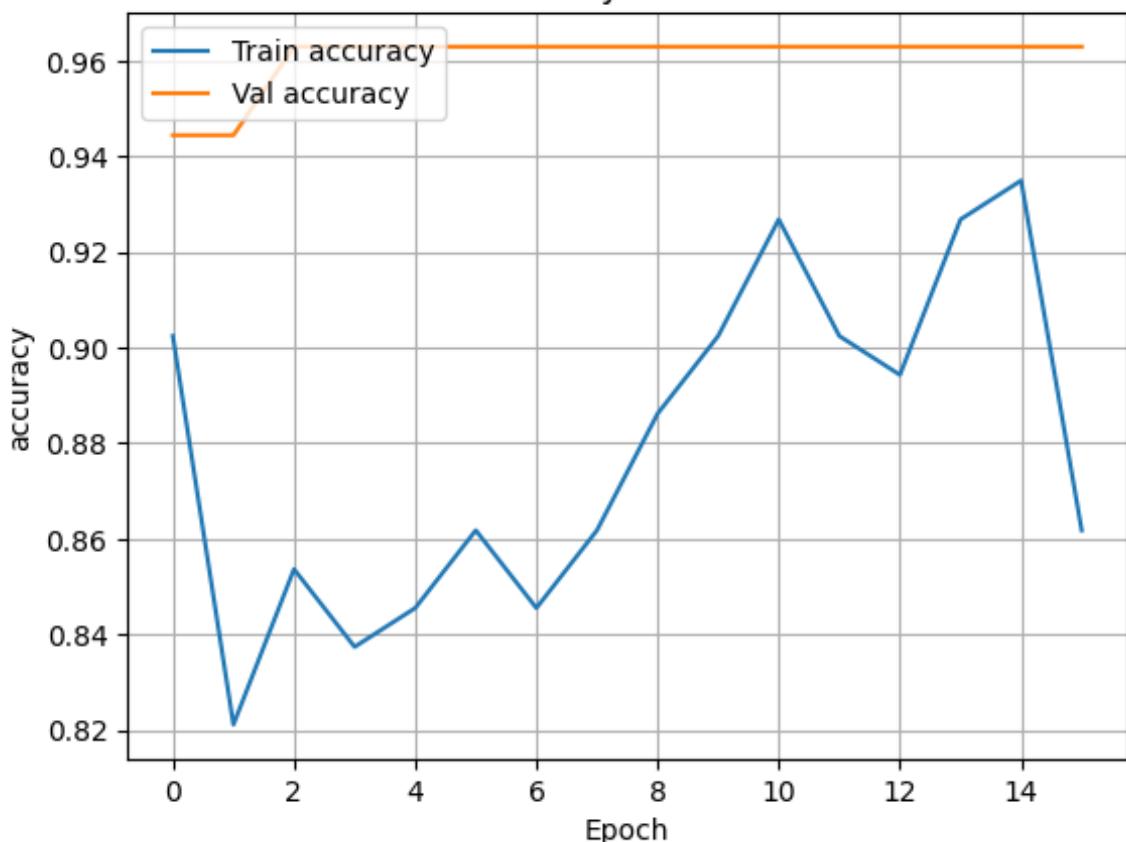
Epoch 15/16
8/8 16s 2s/step - accuracy: 0.9350 - loss: 0.32
97 - val_accuracy: 0.9630 - val_loss: 0.3046

Epoch 16/16
8/8 16s 2s/step - accuracy: 0.8618 - loss: 0.43
39 - val_accuracy: 0.9630 - val_loss: 0.2968

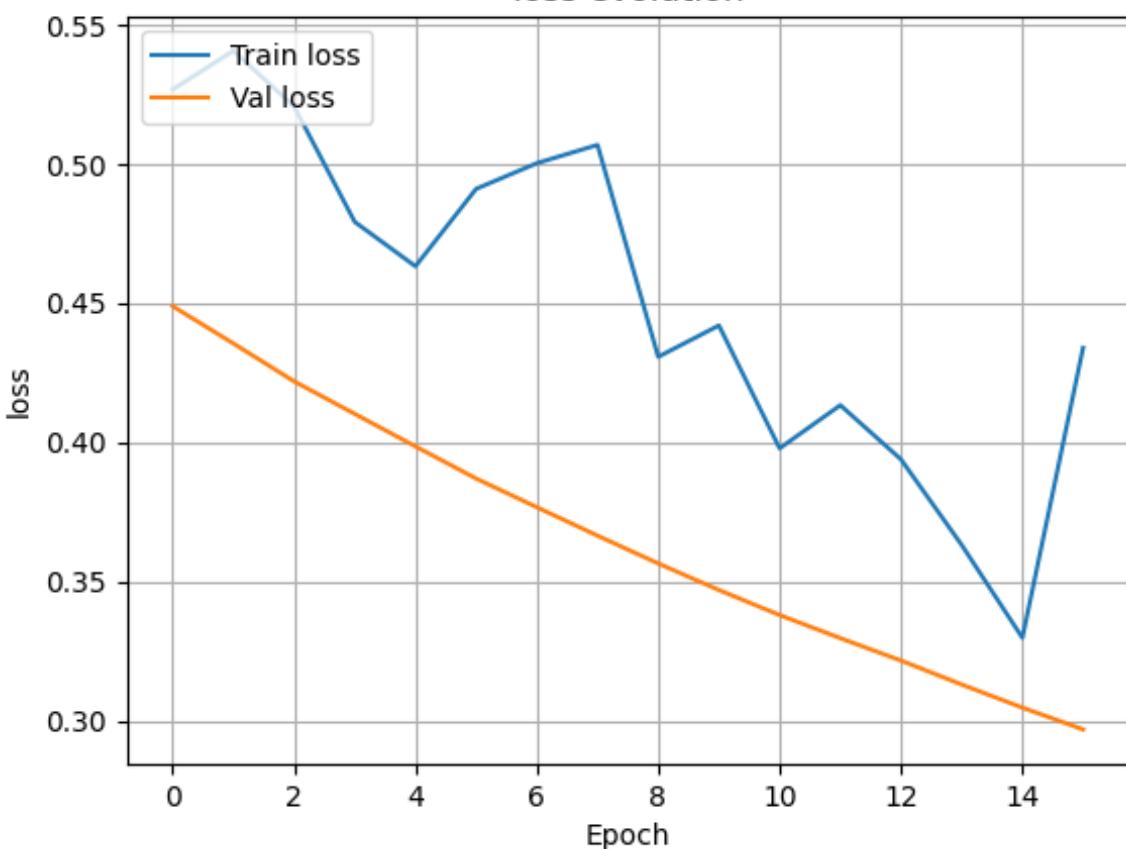
2/2 5s 2s/step - accuracy: 0.9630 - loss: 0.2968

Test loss: 0.29680973291397095, Test accuracy: 0.9629629850387573

accuracy evolution



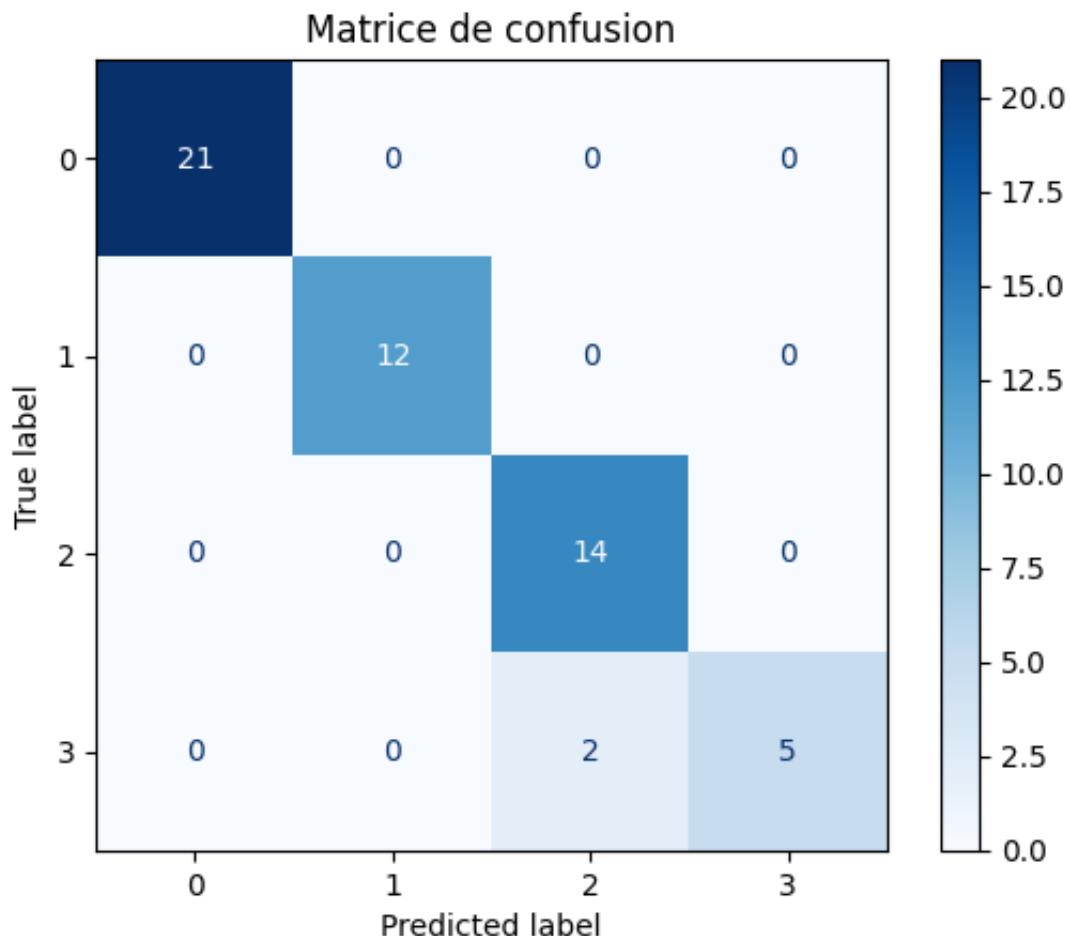
loss evolution



WARNING:tensorflow:5 out of the last 5 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x1208b6d40> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

1/2  **2s** 3s/step WARNING:tensorflow:6 out of the last 6 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x1208b6d40> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

2/2  **5s** 2s/step



```

Classification report:
      precision    recall  f1-score   support

          0       1.00      1.00      1.00      21
          1       1.00      1.00      1.00      12
          2       0.88      1.00      0.93      14
          3       1.00      0.71      0.83       7

   accuracy                           0.96      54
macro avg       0.97      0.93      0.94      54
weighted avg    0.97      0.96      0.96      54

```

Temps d'entraînement : 264.70 secondes

Avec du fine-tuning les performances sont encore améliorées:

```

Classification report:
      precision    recall  f1-score   support

          0       1.00      1.00      1.00      21
          1       1.00      1.00      1.00      12
          2       0.88      1.00      0.93      14
          3       1.00      0.71      0.83       7

   accuracy                           0.96      54
macro avg       0.97      0.93      0.94      54
weighted avg    0.97      0.96      0.96      54

```

Temps d'entraînement : 264.70 secondes

Pour éviter de ruiner les poids appris par VGG, j'ai changé la façon dont je permettais aux couches d'être fine-tuned pour ne laisser que les 4 dernières couches (celles-ci étant les plus spécialisées).

On note aussi que la précision elle même a eu l'air de stagner, car ce fine-tuning n'a corrigé qu'une des 3 erreurs que faisait le modèle auparavant.

Cependant, la loss à bel et bien progressé signifiant sûrement une hausse de la certitude.

```

In [32]: from TP3_utils import load_mnist_with_noise
(X_train, Y_train, X_train_noise), (X_test, Y_test, X_test_noise) = load_mnist_with_noise()

print(X_train_noise.shape, X_train.shape)
print(X_train_noise.min(), X_train_noise.max())
print(X_train.min(), X_train.max())

# Display the train data and a version of it with added noise
for i in range(5):
    plt.subplot(2,5,i+1)
    plt.imshow(X_train[i,:].reshape([28,28]), cmap='gray')
    plt.axis('off')
    plt.subplot(2,5,i+6)
    plt.imshow(X_train_noise[i,:].reshape([28,28]), cmap='gray')
    plt.axis('off')
plt.show()

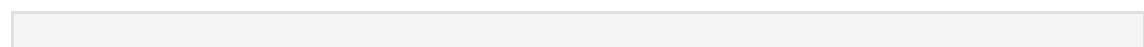
```

(60000, 28, 28, 1) (60000, 28, 28, 1)

0.0 1.0

0.0 1.0





```
In [ ]: from TP3_utils import build_autoencoder

lr_scheduler = ReduceLROnPlateau(monitor='val_loss', factor=0.75, patience=5)
earlyStop = EarlyStopping(monitor='val_loss', min_delta=1e-4, patience=10)

autoencoder = build_autoencoder(X_train[0].shape)
autoencoder.summary()

autoencoder.compile(optimizer=Adam(1e-3), loss='mse')

history_ae = autoencoder.fit(
    X_train_noise,
    X_train,
    validation_data=(X_test_noise, X_test),
    epochs=8,
    batch_size=256,
    callbacks=[earlyStop, lr_scheduler]
)

# Visualisation reconstruction
decoded_imgs = autoencoder.predict(X_test_noise)
for i in range(5):
    plt.subplot(2,5,i+1)
    plt.imshow(X_test[i].reshape(28,28), cmap='gray')
    plt.axis('off')
    plt.subplot(2,5,i+6)
    plt.imshow(decoded_imgs[i].reshape(28,28), cmap='gray')
    plt.axis('off')
plt.show()

affiche(history)
test_loss = autoencoder.evaluate(X_test_noise, X_test)
print(f'Test loss: {test_loss:.4f}')

decoded_imgs = autoencoder.predict(X_test_noise)

for i in range(5):
    plt.subplot(3,5,i+1)
    plt.imshow(X_test_noise[i].reshape(28,28), cmap='gray')
    plt.axis('off')
    if i == 0: plt.ylabel('Bruitée')

    plt.subplot(3,5,i+6)
    plt.imshow(decoded_imgs[i].reshape(28,28), cmap='gray')
    plt.axis('off')
    if i == 0: plt.ylabel('Reconstruite')

    plt.subplot(3,5,i+11)
    plt.imshow(X_test[i].reshape(28,28), cmap='gray')
    plt.axis('off')
```

```

    if i == 0: plt.ylabel('Originale')
plt.show()

print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))

```

Model: "functional_32"

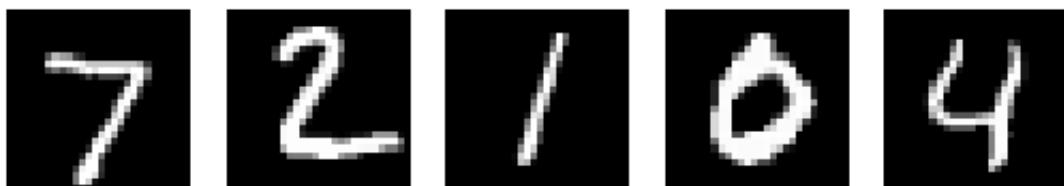
Layer (type)	Output Shape
input_layer_44 (InputLayer)	(None, 28, 28, 1)
conv2d_98 (Conv2D)	(None, 28, 28, 16)
max_pooling2d_48 (MaxPooling2D)	(None, 14, 14, 16)
conv2d_99 (Conv2D)	(None, 14, 14, 32)
max_pooling2d_49 (MaxPooling2D)	(None, 7, 7, 32)
conv2d_100 (Conv2D)	(None, 7, 7, 64)
conv2d_transpose_46 (Conv2DTranspose)	(None, 14, 14, 32)
conv2d_transpose_47 (Conv2DTranspose)	(None, 28, 28, 16)
conv2d_101 (Conv2D)	(None, 28, 28, 1)

Total params: 46,529 (181.75 KB)

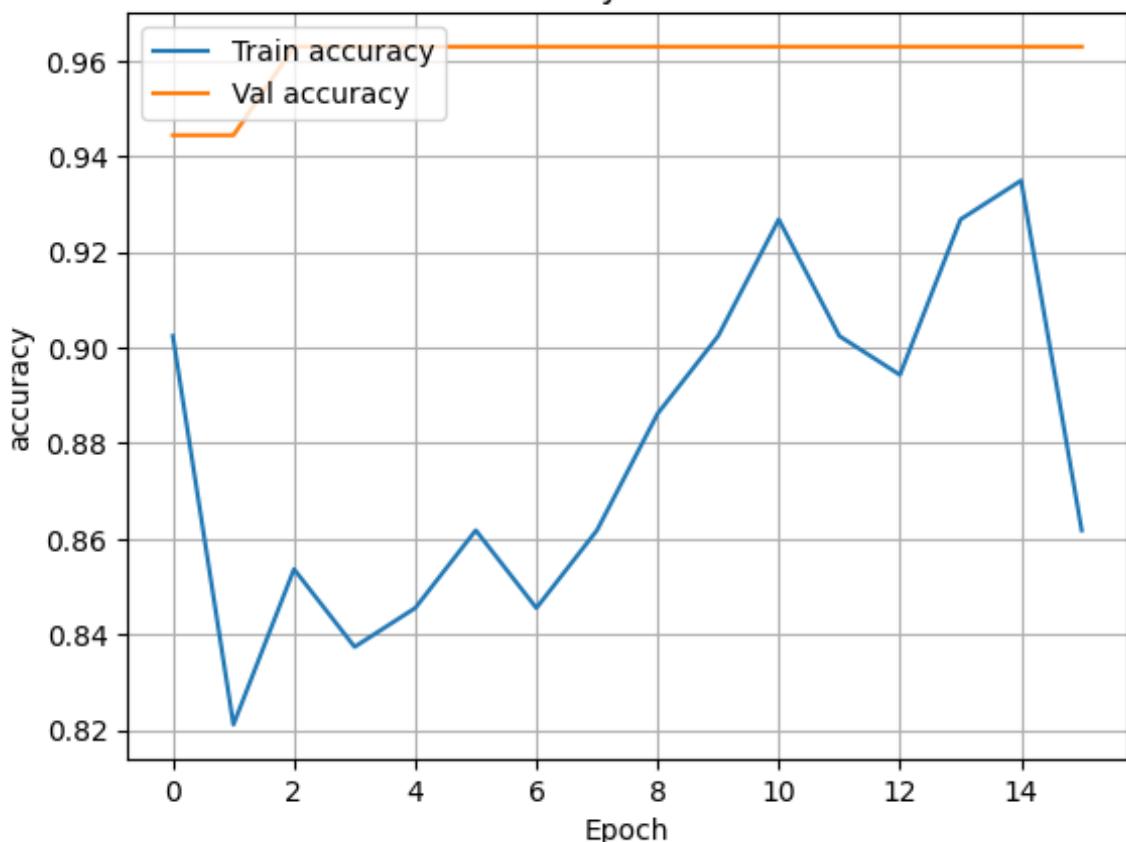
Trainable params: 46,529 (181.75 KB)

Non-trainable params: 0 (0.00 B)

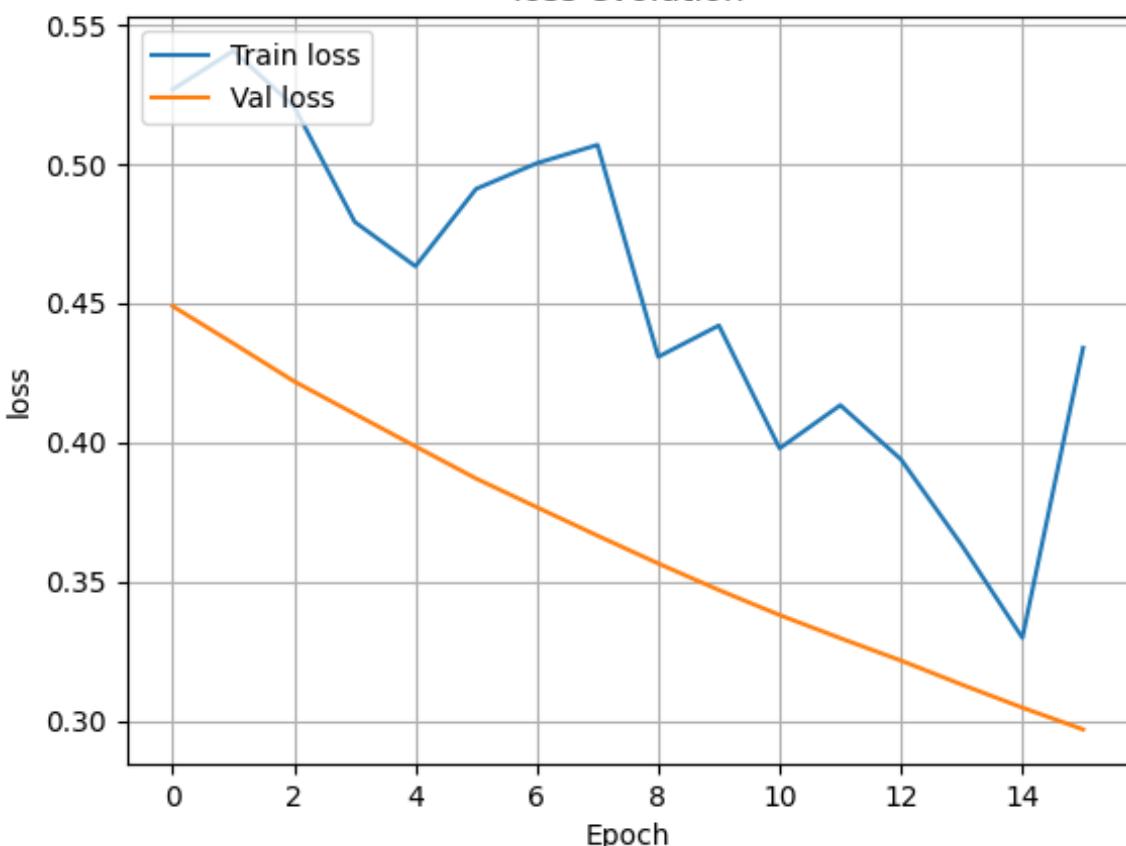
```
Epoch 1/8
235/235 12s 46ms/step - loss: 0.0241 - val_loss: 0.0157 - learning_rate: 0.0010
Epoch 2/8
235/235 10s 44ms/step - loss: 0.0069 - val_loss: 0.0143 - learning_rate: 0.0010
Epoch 3/8
235/235 10s 44ms/step - loss: 0.0058 - val_loss: 0.0136 - learning_rate: 0.0010
Epoch 4/8
235/235 11s 45ms/step - loss: 0.0053 - val_loss: 0.0132 - learning_rate: 0.0010
Epoch 5/8
235/235 10s 44ms/step - loss: 0.0049 - val_loss: 0.0130 - learning_rate: 0.0010
Epoch 6/8
235/235 10s 44ms/step - loss: 0.0047 - val_loss: 0.0130 - learning_rate: 0.0010
Epoch 7/8
235/235 0s 42ms/step - loss: 0.0045
Epoch 7: ReduceLROnPlateau reducing learning rate to 0.0007500000356
230885.
235/235 10s 44ms/step - loss: 0.0045 - val_loss: 0.0129 - learning_rate: 0.0010
Epoch 8/8
235/235 10s 44ms/step - loss: 0.0043 - val_loss: 0.0129 - learning_rate: 7.5000e-04
Restoring model weights from the end of the best epoch: 5.
313/313 1s 2ms/step
```



accuracy evolution



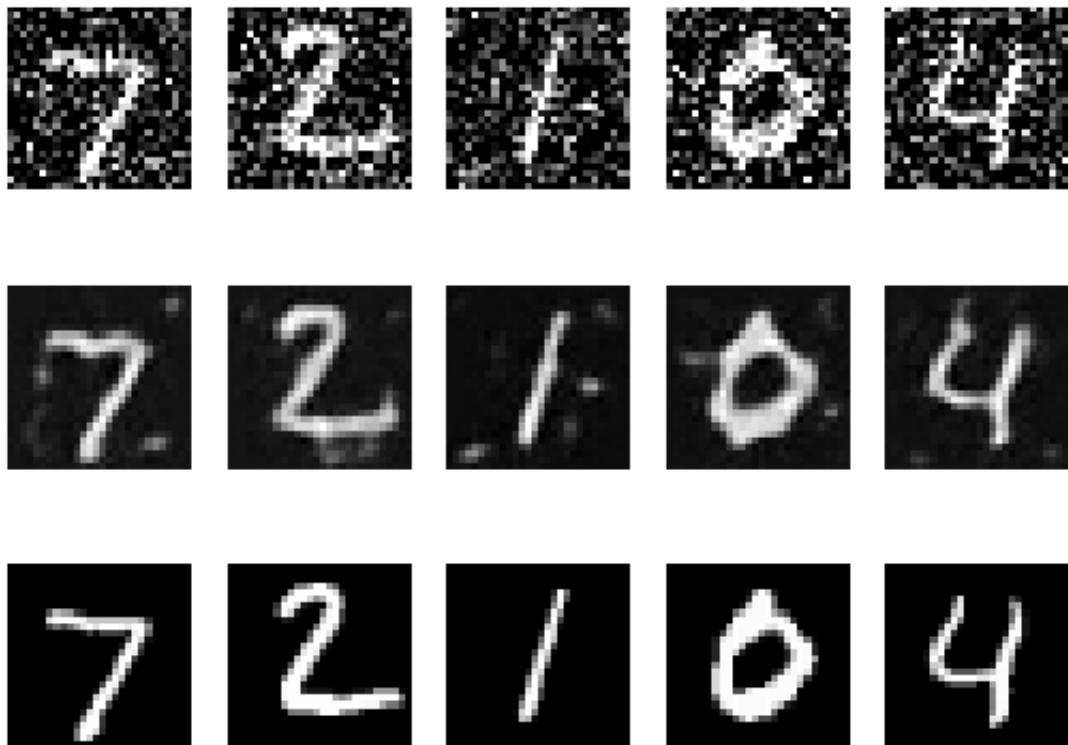
loss evolution



313/313 ━━━━━━ 1s 2ms/step - loss: 0.0130

Test loss: 0.0130

313/313 ━━━━━━ 1s 2ms/step



Temps d'entraînement : 264.70 secondes

```
In [ ]: from TP3_utils import build_vae
lr_scheduler = ReduceLROnPlateau(monitor='val_loss', factor=0.8, patience=5)
earlyStop = EarlyStopping(monitor='val_loss', min_delta=1e-4, patience=10)

latent_dim = 2
coeff = 0.01
vae, encoder, decoder = build_vae(latent_dim=latent_dim, coeff=coeff)
vae.compile(optimizer=Adam(1e-3), loss='mse')

history = vae.fit(
    X_train_noise, X_train,
    validation_data=(X_test_noise, X_test),
    epochs=20,
    batch_size=128,
    callbacks=[earlyStop, lr_scheduler]
)
```

```
Epoch 1/20
469/469 ━━━━━━━━━━ 14s 28ms/step - loss: 0.0781 - val_loss: 0.0680 - learning_rate: 0.0010
Epoch 2/20
467/469 ━━━━━━ 0s 25ms/step - loss: 0.0675
Epoch 2: ReduceLROnPlateau reducing learning rate to 0.0007500000356230885.
469/469 ━━━━━━━━━━ 12s 26ms/step - loss: 0.0675 - val_loss: 0.0676 - learning_rate: 0.0010
Epoch 3/20
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0674 - val_loss: 0.0675 - learning_rate: 7.5000e-04
Epoch 4/20
468/469 ━━━━━━ 0s 26ms/step - loss: 0.0672
Epoch 4: ReduceLROnPlateau reducing learning rate to 0.0005625000048894435.
```

```
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 7.5000e-04
Epoch 5/20
469/469 ━━━━━━━━━━ 13s 28ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 5.6250e-04
Epoch 6/20
467/469 ━━━━━━ 0s 27ms/step - loss: 0.0672
Epoch 6: ReduceLROnPlateau reducing learning rate to 0.0004218749818392098.
469/469 ━━━━━━━━━━ 13s 28ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 5.6250e-04
Epoch 7/20
469/469 ━━━━━━━━━━ 13s 28ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 4.2187e-04
Epoch 8/20
468/469 ━━━━━━ 0s 26ms/step - loss: 0.0672
Epoch 8: ReduceLROnPlateau reducing learning rate to 0.00031640623637940735.
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 4.2187e-04
Epoch 9/20
469/469 ━━━━━━━━━━ 13s 28ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 3.1641e-04
Epoch 10/20
467/469 ━━━━━━ 0s 26ms/step - loss: 0.0672
Epoch 10: ReduceLROnPlateau reducing learning rate to 0.00023730468819849193.
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 3.1641e-04
Epoch 11/20
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 2.3730e-04
Epoch 12/20
467/469 ━━━━━━ 0s 26ms/step - loss: 0.0672
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.00017797851614886895.
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 2.3730e-04
Epoch 13/20
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 1.7798e-04
Epoch 14/20
468/469 ━━━━━━ 0s 26ms/step - loss: 0.0672
Epoch 14: ReduceLROnPlateau reducing learning rate to 0.0001334838816546835.
469/469 ━━━━━━━━━━ 13s 27ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 1.7798e-04
Epoch 15/20
469/469 ━━━━━━━━━━ 13s 28ms/step - loss: 0.0673 - val_loss: 0.0675 - learning_rate: 1.3348e-04
Epoch 15: early stopping
Restoring model weights from the end of the best epoch: 1.
```

```
In [53]: decoded_imgs = vae.predict(X_test_noise)
for i in range(5):
    plt.subplot(2,5,i+1)
```

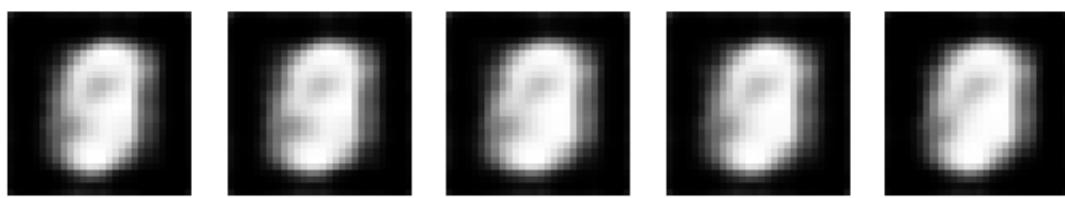
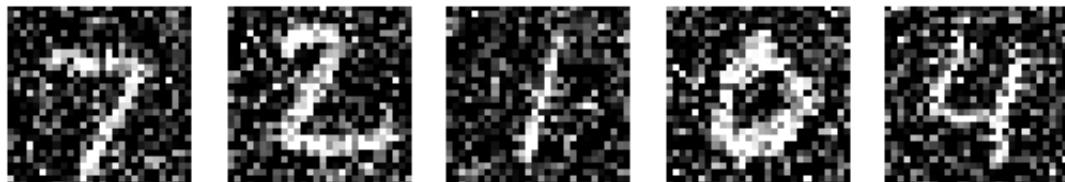
```

plt.imshow(X_test_noise[i].reshape(28,28), cmap='gray')
plt.axis('off')
if i==0: plt.ylabel('Bruitée')

plt.subplot(2,5,i+6)
plt.imshow(decoded_imgs[i].reshape(28,28), cmap='gray')
plt.axis('off')
if i==0: plt.ylabel('Reconstruite')
plt.show()

```

313/313 ━━━━━━ 1s 3ms/step



In [54]:

```

n = 10
grid_x = np.linspace(-3, 3, n)
grid_y = np.linspace(-3, 3, n)

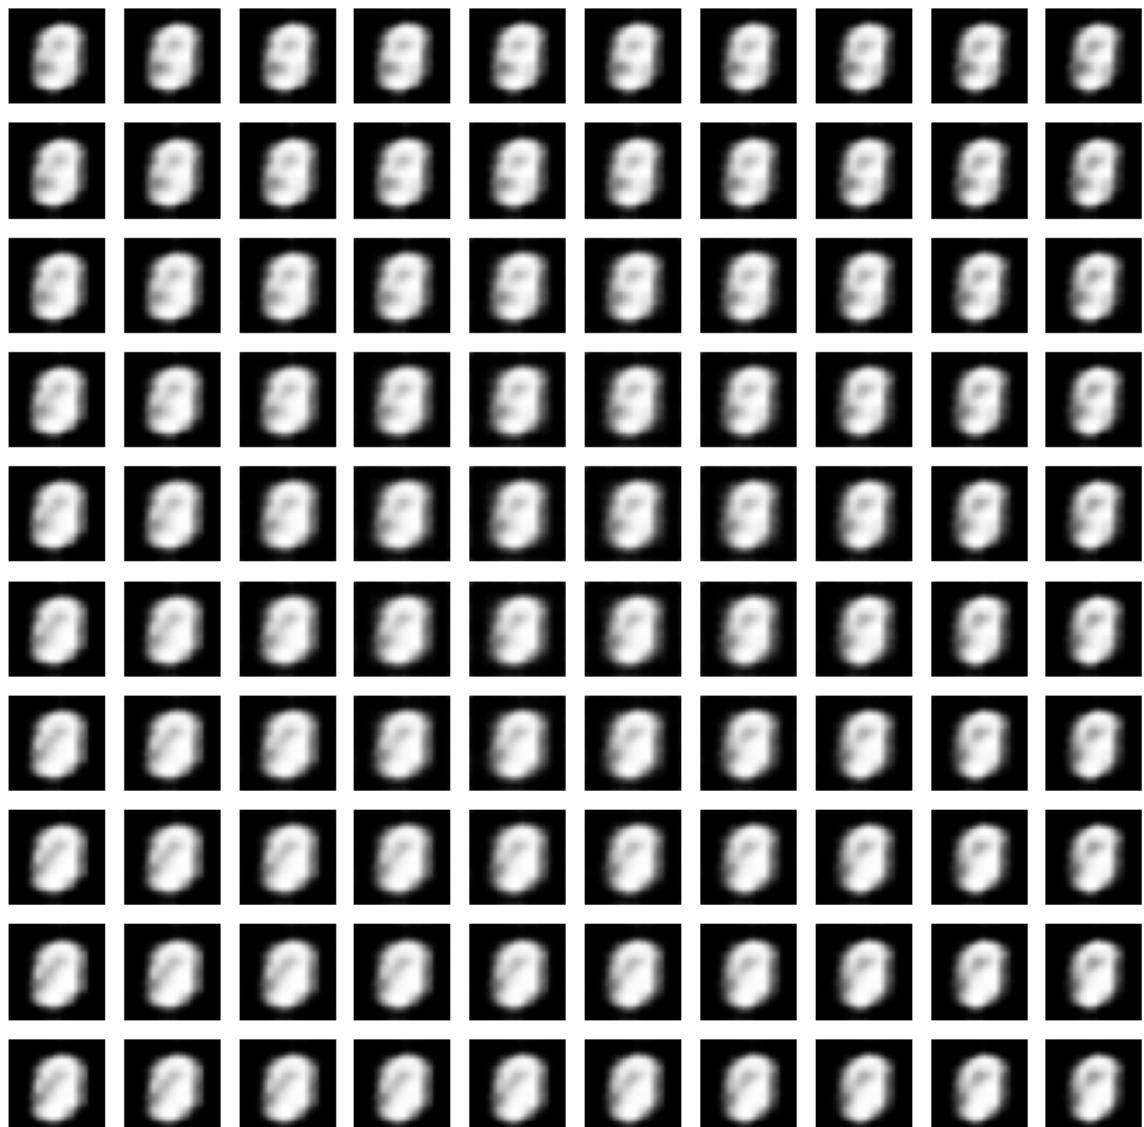
plt.figure(figsize=(10,10))
for i, xi in enumerate(grid_x):
    for j, yi in enumerate(grid_y):
        z_sample = np.array([[xi, yi]])
        x_decoded = decoder.predict(z_sample)
        plt.subplot(n, n, i*n+j+1)
        plt.imshow(x_decoded[0].reshape(28,28), cmap='gray')
        plt.axis('off')
plt.suptitle("Espace latent du VAE")
plt.show()

```

1/1 ━━━━━━ 0s 29ms/step
1/1 ━━━━━━ 0s 11ms/step
1/1 ━━━━━━ 0s 11ms/step
1/1 ━━━━━━ 0s 10ms/step
1/1 ━━━━━━ 0s 10ms/step
1/1 ━━━━━━ 0s 10ms/step
1/1 ━━━━━━ 0s 11ms/step
1/1 ━━━━━━ 0s 10ms/step
1/1 ━━━━━━ 0s 10ms/step
1/1 ━━━━━━ 0s 11ms/step
1/1 ━━━━━━ 0s 10ms/step
1/1 ━━━━━━ 0s 10ms/step

1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 9ms/step
1/1 0s 10ms/step
1/1 0s 9ms/step
1/1 0s 10ms/step
1/1 0s 9ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step
1/1 0s 10ms/step

Espace latent du VAE



In []: