

CNN Avec peu de données

Chargement du code

```
In [ ]: %load_ext autoreload  
%autoreload 2  
from TP3_utils import load_dataset  
import numpy as np  
from matplotlib import pyplot as plt  
  
print("un")  
DATA_PATH = 'Data'  
CATEGORIES = ["accordion", "anchor", "barrel", "binocular"]  
print("deux")  
# Chargement du dataset  
(X_train, Y_train), (X_test, Y_test), num_classes = load_dataset(DATA_PATH)  
print("trois")  
  
# Vérification rapide  
data_shape = X_train[0].shape  
print("Training data shape:", data_shape)  
print("Training labels shape:", Y_train.shape)  
print("Test data shape:", X_test.shape)  
print("Test labels shape:", Y_test.shape)  
print("Number of classes:", num_classes)
```

Redimensionnement sur les images de X_train
pour y_train on passe du label brut au vecteur one-hot (classification multiclass)

```
In [ ]: import time  
import tensorflow as tf  
from tensorflow import keras  
from keras.optimizers import Adam  
from TP3_utils import CNN, affiche, eval_classif  
# from keras.callbacks import ReduceLROnPlateau, EarlyStopping  
  
# lr_schedulerBis = ReduceLROnPlateau(monitor='val_loss', factor=0.9, patience=3)  
# earlStopBis = EarlyStopping(monitor='val_loss', min_delta=1e-3, patience=3)  
  
lr=1e-4  
batch_size=min([X_train.shape[0], 256])  
epochs=16  
ad= Adam(learning_rate=lr)  
  
model = CNN(data_shape)  
model.summary()  
  
model.compile(  
    loss='categorical_crossentropy',  
    optimizer=ad,
```

```

        metrics=['accuracy']
    )

tps1 = time.time()
history =model.fit(
    X_train,
    Y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(X_test, Y_test)
#    callbacks=[earlyStopBis, lr_schedulerBis],
)
tps2 = time.time()

# Evaluation
loss, accuracy = model.evaluate(X_test, Y_test)
print(f'Test loss: {loss}, Test accuracy: {accuracy}')

affiche(history)
preds = model.predict(X_test)
eval_classif(Y_test, preds)
print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))

```

Premier essaie:

- Clairement overfitting: loss value stagne vite, l'accuracy est de 1 en train, mais .85 en test

```
In [ ]: from keras.layers import Input, Flatten, Dense
from TP3_utils import MLP_transfer

input_tensor = Input(shape=(224,224,3))
VGG = tf.keras.applications.VGG16(weights='imagenet',
    include_top=True,
    input_tensor=input_tensor
)
VGG.summary()

feature_train = VGG.predict(X_train)
feature_test = VGG.predict(X_test)
```

```
In [ ]: # Transfer learning
input_tensor = Input(shape=(224,224,3))
VGG = tf.keras.applications.VGG16(weights='imagenet',
    include_top=False,
    input_tensor=input_tensor
)
for layer in VGG.layers:
    layer.trainable = False

model_transfer = keras.Sequential()
model_transfer.add(VGG)
model_transfer.add(MLP_transfer(VGG.output_shape[1:], num_classes))
model_transfer.summary()
```

```

model_transfer.compile(
    loss='categorical_crossentropy',
    optimizer=ad,
    metrics=['accuracy']
)

tps1 = time.time()
history =model_transfer.fit(
    X_train,
    Y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(X_test, Y_test)
#    callbacks=[earlyStopBis, lr_schedulerBis],
)
tps2 = time.time()
loss, accuracy = model_transfer.evaluate(X_test, Y_test)
print(f'Test loss: {loss}, Test accuracy: {accuracy}')

affiche(history)
preds = model_transfer.predict(X_test)
eval_classif(Y_test, preds)
print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))

```

In []: # Fine-tuning

```

In [ ]: from tensorflow.keras.datasets.mnist import load_data

(X_train, y_train), (X_test, y_test) = load_data()
# Preprocess input data
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_train.shape[1], X_train.shape[2])
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255
X_test /= 255

X_train_noise = X_train + 0.2 * np.random.normal(loc=0.0, scale=1.0)
X_test_noise = X_test + 0.4 * np.random.normal(loc=0.0, scale=1.0)

X_train_noise = np.clip(X_train_noise, 0.0, 1.0)
X_test_noise = np.clip(X_test_noise, 0.0, 1.0)
# Display the train data and a version of it with added noise
for i in range(5):
    plt.subplot(2,5,i+1)
    plt.imshow(X_train[i,:].reshape([28,28]), cmap='gray')
    plt.axis('off')
    plt.subplot(2,5,i+6)
    plt.imshow(X_train_noise[i,:].reshape([28,28]), cmap='gray')
    plt.axis('off')
    plt.show()

```

In []: from TP3_utils import auto_encoder

```
model_ae = auto_encoder(data_shape)

model_ae.compile(
    loss="mean_square_error",
    optimizer=ad,
    metrics=['accuracy']
)

history =model_ae.fit(
    X_train,
    Y_train,
    batch_size=32,
    epochs=20,
    verbose=1,
    validation_data=(X_test, Y_test)
#    callbacks=[earlyStopBis, lr_schedulerBis],
)

loss, accuracy = model_ae.evaluate(X_test, Y_test)
print(f'Test loss: {loss}, Test accuracy: {accuracy}')

affiche(history)
preds = model_ae.predict(X_test)
eval_classif(Y_test, preds)
print("Temps d'entraînement : {:.2f} secondes".format(tps2 - tps1))
```